Same Samples: Porting a Block-Based Synthesizer Firmware from C++ to Faust, Checked Against the Firmware Itself

Stéphane Letz
GRAME-CNCM, Lyon
letz@grame.fr

October 2026

Abstract

Mutable Instruments Elements is a Eurorack modal synthesis voice whose firmware is open source: three exciters (a bow, a blow with a reed tube, a strike that morphs from recorded mallets to bouncing particles), a resonator of 52 zero-delay-feedback band-pass modes with 8 bowed waveguides, a single string or a chord of five dispersive strings, and a reverb whose delay memory is 16 bits, in about 3,700 lines of C++ run at 32 kHz in blocks of 16 samples. We port all of it to one Faust library, compiled by the Rust Faust compiler faust-rs, and ask for the same samples as the firmware, not the same kind of sound. Five things make a block-based engine with a shared random generator into a sample stream that gives its numbers: block-rate work in `ondemand` domains clocked by the block start, ramps that accumulate as the C++ accumulates, `ondemand` domains that do not fire wherever the C++ skips a computation, so that frozen filters and delay lines resume where they were, delay lines read with the C++'s ring-buffer semantics, and the generator's words handed to each consumer by its k-step jump in the C++'s order. Against the firmware's own `Part` and `Voice`, compiled unchanged, on 11 scenarios that exercise every exciter model, the three resonators, moving parameters, an external input and the reverb, the Faust model agrees to -72 dB of the peak or better in single precision, -100 dB on most; in double precision it drifts where the firmware accumulates phases in float32. The port found one defect of faust-rs: a delay of unbounded interval is sized to 2^{31} samples where the C++ compiler refuses the program.

1 Introduction

A port of a synthesizer between languages is usually judged by ear. That judgement is weak where the model is noisy, nonlinear or chaotic: a bowed string, a granular noise or bouncing particles never sound twice the same, and two renders that sound alike can still disagree by several decibels in a register nobody listened to. We take the stance of an earlier port of a physically modelled piano [10]: the C++ is its own reference. A small program, `ref/dump.cpp`, links the firmware's DSP sources [5] unchanged, dumps every table the Faust model needs, and renders scenarios through the firmware's `elements::Part` (the whole module) and `elements::Voice` (the voice alone). The Faust model [12] is rendered by `faustprobe`, the offline probe of faust-rs [7], under the same control changes at the same samples, and the two are compared sample by sample.

The piano was a model that prepares its notes offline, and its port was a reformulation. Elements is the opposite case: a real-time engine written for a Cortex-M4, whose every sample is already computed in a loop. The difficulty is not to reformulate but to reproduce exactly an engine organized around blocks of 16 samples, with control-rate work at the top of each block, per-block ramps, components that are simply not called when unused, and one random generator shared by four consumers whose draws interleave block by block. Faust has no blocks: a program computes all its signals at every sample. This paper describes how the block structure was written into a sample stream (Sections 3–5), what the comparison measures (Section 6), and what the compiler made easy and hard (Section 7).

2 The firmware

The DSP of Elements [5] (Figure 1) runs at 32 kHz, the module's codec rate, in blocks of 16 samples (`kMaxBlockSize`). `Part::Process` reads the patch and the performance state (gate, note, strength, modulation) once per block, renders the voice, mixes it to stereo, soft-limits each channel and applies the reverb.

Exciters. The *bow* is a value that flips sign when a uniform draw falls under a threshold, mixed with the draw and low-passed by a state-variable filter [15]; it also sets the bow pressure of the bowed modes. The *blow* reads a 40,963-sample noise recording at a rate set by its timbre and jumps back to a restart point when a draw falls under 1 %; above two thirds of its level a reed and bore model in the spirit of McIntyre,

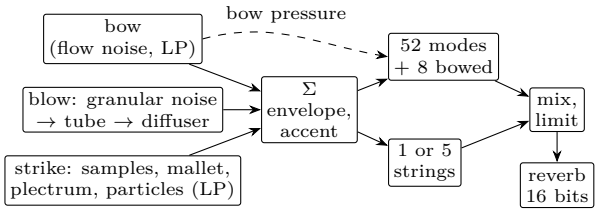


Figure 1: The voice and the space section of Elements. The raw excitation, the resonator's center and its side are the voice's three outputs; `Part` mixes them by the *space* control.

Schumacher and Woodhouse [11] joins in; four all-passes diffuse the result. The *strike* morphs with its meta control between nine mallet and hit recordings (128,013 samples in all, crossfaded two by two), a filtered pulse, a plectrum (a small pulse, then the main one after a delay, and a damping request that mutes the resonator on release) and bouncing particles whose intervals and gains follow random draws. An envelope with four contours (percussive, AD, ADSR, AR) and an accent table driven by the strength input scale them.

Resonator. The modal resonator [1] is 52 band-pass state-variable filters at harmonics of the note stretched by a stiffness that the *geometry* control sets (compressed partials, harmonic, then bars, plates and bells), with Q falling from mode to mode according to *damping* and *brightness*. The modes are summed with the amplitudes of a cosine comb, the frequency-domain image of a pickup at *position*, into the center output, and with those of a second pickup moved by a slow triangle LFO into the side output. The first 8 modes are doubled by bowed waveguides [4]: a delay line of one period in a loop with a normalized band-pass, all excited by the bow table [2] of their summed output. Hidden modes replace the resonator by one string or by five strings tuned to a chord: Karplus–Strong strings [8, 9, 14] with Hermite-interpolated delay reads, an FIR and an SVF for the damping, dispersion through an all-pass, noise on the delay and a "curved bridge" nonlinearity, and a comb pickup; below 15.7 Hz, where the 2048-sample line is too short, a string runs slower than the sample rate and is up-sampled by linear interpolation.

Space. The *space* control first feeds the raw exciter to the left output, then widens the stereo image, then raises a Griesinger reverberator in the topology described by Dattorro [3], and freezes it at the end of its travel. Its ten delay lines share one 32,768-cell buffer of 16-bit samples, and two approximate cosine LFOs modulate two of its reads.

3 From blocks to a stream

The block clock. The model counts samples from the first one it computes, as the firmware's first block starts at its first sample:

```
bpos = ba.time % BS;          // BS = 16
blk  = bpos == 0;             // the first sample of a block
_hb(x) = (blk, x) : ondemand(_);  // a control held for its
      block
```

In faust-rs, `ondemand(f)` runs `f` only on the samples where its first input is non-zero, holds its outputs otherwise, and advances its state (recursions, delays) once per firing. Everything the C++ computes once per block is such a domain clocked by `blk`: the gate flags, the envelope (a five-state machine stepping one phase increment per block), the strike exciter, the chord choice with its hysteresis, the strings' parameters, and `Resonator::ComputeFilters`, which refreshes the first 25 modes every block and the 27 others on alternate blocks by the parity of its own call count. The controls are held per block: one moved in the middle of a block acts at the next, as on the module.

Ramps. The C++ moves a few values across a block by adding, at every sample, an increment computed at the block start, `(target - value) / 16`: the envelope, the strength, the pickup position, four string parameters. The sum of 16 increments does not land exactly on the target in float32, and the C++ carries the difference into the next block. `_ramp` does the same, from the value reached at the end of the previous block:

```
_ramp(clk, target, x0) = (\(v, inc).(vn, incn with {
    vp = _init(x0, v);
    incn = _if(clk, (target - vp) / float(BS), inc);
    vn = vp + incn; })) ~ (_, _) : _, !;
```

What the C++ does not process. In string mode the modal resonator is not called; neither are the bowed modes above the number of active modes, the four extra strings in single-string mode, the dispersion all-pass when it is out of the loop, nor a slow string between its ticks. Their filters and delay lines stand still and resume where they were. Each is written as an `ondemand` domain whose clock is that condition; a string, for instance, is a domain of the voice, its tick a domain of the string, and the dispersion all-pass a domain of the tick, three levels deep inside recursions. The modes above the number of active ones only keep their state with a select.

Delay lines as ring buffers. A `DelayLine` read at offset k before its write returns what was written k writes ago, and N writes ago when k is a multiple of the size N; read after its write, offset k returns what was written $k-1$ writes ago, the current value at 1. The model reads a recursion's fed-back value `wp` (what was written one write ago) the same way, `_ring(N, wp, k) = wp @ ((k - 1 + N) % N)`, which also gives the C++'s corner cases (a comb pickup at less than one sample, a read that wraps).

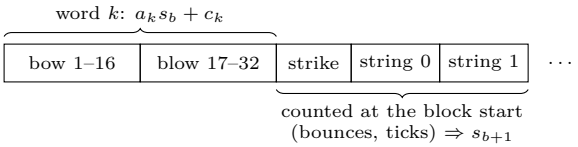


Figure 2: The generator's words within one block, in the order the C++ draws them.

4 One shared generator

`stmlib::Random` is one linear congruential generator, $s \mapsto 1664525\,s + 1013904223 \bmod 2^{32}$ [13], shared by the whole voice. Within a block the bow draws once per sample, then the blow once per sample, then the particles of the strike (one word on the gate's rising edge, two or three per bounce, at most one bounce per block since the next interval is at least 96 samples), then each string once per tick (Figure 2). The consumers do not draw in time order: the blow's first word of a block follows the bow's last, and the strings' words depend on the particles' bounces in the same block.

The model keeps the generator's state s_b at the start of block b and gives every consumer the word it would draw. The state k steps after s is $a_k s + c_k$ with $a_k = a^k$ and $c_k = c \sum_{i<k} a^i$ modulo 2^{32}, tabulated for $k = 0 \ldots 48$ by iterating the C++'s own step; Faust's integers are 32-bit and wrap as the C++'s unsigned arithmetic does. The bow takes word `bpos` + 1 and the blow word `bpos` + 17:

```
bow_sig  = bow_flow(rise & blk, _lcg_jump(bpos + 1, sb), ...);
blow_sig = blow_granular(_lcg_jump(bpos + 17, sb), ...) : tube
     (...) : ... : diffuser;
```

The strike's consumption depends on its own state, so the strike is computed at block rate: `random_block`, clocked by `blk`, unrolls the 16 steps of each strike model from word 33 and returns the 16 samples, the damping they request, the strike's state and the generator's state after its draws. Each string then starts after the words of the strings before it; a string below 15.7 Hz ticks fewer than 16 times per block, so its ticks are counted at the block start by running its resampling phase through the block, the same floating-point operations the string performs sample by sample. The state of the next block follows. Unsigned comparisons flip the sign bit of both operands; unsigned shifts mask; the word becomes a float in $[0, 1)$ as `GetFloat` makes it.

5 The reverb in 16 bits

`FxEngine` stores each delay line in 16 bits: a write truncates $32768\,x$ to an integer and saturates it, a read divides by 32768. The model quantizes every write the same way (`_q16`) and keeps the unquantized accumulator within the sample, as the C++ does. The first input all-pass is "smeared": its line is read at an LFO-modulated offset between 10 and 90 samples, and the interpolated value is written back at offset

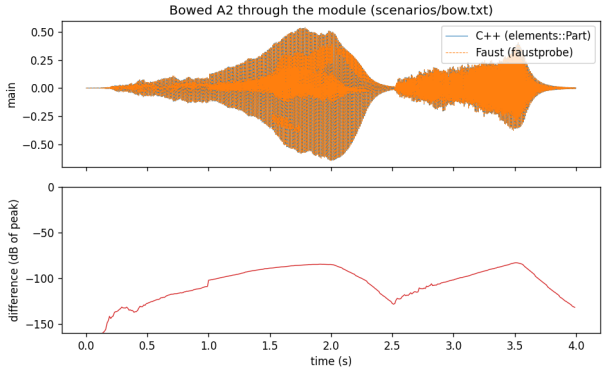


Figure 3: The bow scenario through the module, C++ and Faust superimposed (top) and their difference (bottom), single precision. The bowed modes are a nonlinear loop that amplifies the last-place differences; the worst is −83 dB of the peak.

100 of the same line, where the all-pass reads its tail 49 samples later. In the model this is a delay of a quantized interpolation:

```
tail1 = _q16(sa + (sb - sa) * o1f) @ 49;   // sa, sb: the line
      at o1i, o1i + 1
```

The two LFOs are approximate cosine oscillators stepped every 32 samples, recursions of coefficient $c = 1.99999702$ for the second one. With c that close to 2, one unit in the last place of a float32 changes the oscillator's frequency by about 2 %: the coefficient computed from the C++'s expression in double precision moved the reverb's read offsets enough that the space scenario agreed to −26 dB in single precision; the float32 value, dumped from the C++, brought it to −74 dB. Such constants (two LFO coefficients, the particles' probabilities, the strings' initial delay) are read back from the compiled firmware with the tables.

6 Results

Method. Each scenario is a list of control changes at block boundaries (`FRAME NAME VALUE`). The reference program renders it through the firmware's `Part` or `Voice`, compiled without fused multiply-add so that the compiler cannot change the reference's arithmetic; `faustprobe` renders the Faust module or voice at 32 kHz under the same changes, in single and in double precision. The difference is the largest absolute difference over the whole render, in dB of the C++ channel's peak. Table 1 gives the worst channel of each of the 44 renders (`make results`, Apple M5, Apple clang 21, Elements `f4cfa76`, faust-rs `172305ea`).

Single precision. The model performs the C++'s float operations and should give its float results; it does, to within the reordering of sums and products that the Faust compiler's normal form is free to make, which floating-point arithmetic does not forgive [6]. Where a difference of one unit in the last place crosses a threshold it can grow: the reverb rounds a value to the other 16-bit step and the difference circulates in

		single precision		double precision	
scenario	what it exercises	module	voice	module	voice
strike	the default patch, a mallet	−103	−90	−123	−111
samples	the strike's sample player, three samples	−98	−82	−52	−35
plectrum	plectrum and palm mute	−116	−97	−67	−52
particles	random bounces, three gates	−105	−97	−38	−28
bow	bow, ADSR, moving position and timbre	−83	−84	−58	−35
blow	granular noise, then the reed tube	−105	−95	−56	−46
string	one string, dispersion, endless damping	−101	−101	−58	−60
chords	five strings, a string below 15.7 Hz, bowing	−102	−103	−88	−92
moves	moving parameters, model switched under a note	−84	−85	−54	−42
space	raw, stereo, reverb, frozen reverb	−72	−97	−50	−39
input	an external signal into the resonator	−105	−86	−95	−76

Table 1: Largest difference between the Faust model and the firmware over each render (2 to 6 s), worst channel, in dB of that channel's peak. Module: `elements::Part` (main, aux); voice: `elements::Voice` (raw, center, side). `results/compare.txt` has every channel.

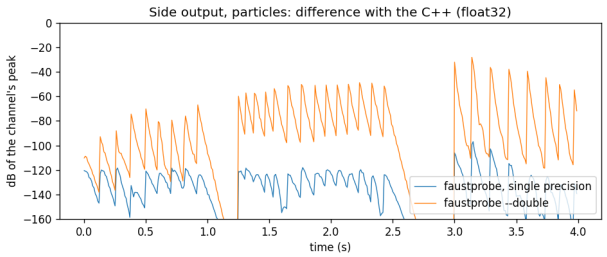


Figure 4: The side output of the particles scenario: difference with the firmware in single and double precision. The double-precision model drifts from the C++'s float32 LFO phase; the single-precision one does not.

its loops (space, −72 dB of the peak, −77 dBFS, a few units of the 16-bit memory); the bowed modes amplify it (bow, −83 dB, Figure 3). The random paths, where one flipped comparison would desynchronize a whole note (particles, blow), agree to −95 dB and better.

Double precision. The firmware accumulates a few phases in float32: the side pickup's LFO phase, the strings' delay ramps, the reverb's LFOs. A double-precision model accumulates the same increments more exactly and drifts away from the firmware (Figure 4). The raw excitation, which accumulates nothing for long, agrees to −99 dB or better in every scenario; the side output and the strings carry the drift.

Invariances. The module gives the same samples at host block sizes of 1, 7 and 512 frames, after a reset, and from a second compilation (`faustprobe --check all`). The comparison was made to fail once: changing the resonator's input gain from 0.125 to 0.126 gives −42 dB on the center output and a failing exit status.

Cost. On 100 s of a held bowed note (`make cost`), the firmware's C++ costs 3.4, 1.5 and 3.9 ms of compute per second of audio with the modal resonator, one string and five strings. The Faust module compiled by faust-rs to C++ and by clang at `-O3` costs 10.7, 2.8 and 4.9 ms, and 17.1, 6.7 and 13.7 ms in `faustprobe`'s Cranelift JIT. The difference lies in the modes: with 26 modes instead of 52 the native module costs 6.4 ms instead of 10.4, so the 52 modes are about 8 ms of its 10.7. The generated code keeps each mode's state in the DSP structure and shifts it every sample where the C++ walks an array of filters; whether that is the whole cause was not measured. Compiled to WebAssembly, one voice and the body cost 13.6 to 16.9 ms per second under node, and agree with `faustprobe` to −78 dB or better; the web version of the repository plays up to eight voices in an AudioWorklet at 32 kHz.

7 The compiler

What was there. The port needed `ondemand` nested three deep inside recursions, 32-bit integer arithmetic that wraps, bit operations for the unsigned comparisons, and `waveform` tables of 128,013 samples; all worked as documented, and the voice compiles in 1.2 s. The block-rate strike, written as recursive definitions unrolled 16 times over tuples of state, compiled without the exponential evaluation that the piano port had met [10], since fixed in faust-rs.

A delay of unbounded interval. The reverb's modulated reads were first written as `wp @ (int(10 + 80 * lfo) - 1)`. The LFO is a recursion, whose interval the analysis cannot bound. The C++ compiler refuses such a program and gives the interval in its message ("possible negative values ... used in delay expression"). faust-rs accepts it and sizes the line to 2^{31} samples: the generated class is 8 GiB, a C++ host that allocates it on the stack fails at once, and `faustprobe` renders the program normally, which hid the problem until the native build. The reads are now clamped to the offsets the C++ reads (10 to 90, 6211 to 6311), which changes nothing in the output and bounds the lines for both compilers; the reduced case is in `faust-rs-issues/unbounded-delay`.

A comma. `_if(c, x : _, !, y)` passes four arguments, not three: in an argument list the comma separates arguments before the composition binds. The first complete voice, with a dozen such calls, ran for more than two minutes without an answer; its components compiled one by one (`faustprobe --eval`) reported the arity errors in a tenth of a second each. Every projection inside an argument is now parenthesized.

8 What is not ported, and conclusion

The module's easter egg (another synthesis engine), the panic reset of `Part` (when the resonator level exceeds 200, which no scenario reaches), and the hardware around `Part` (the pots' smoothing and calibration, the input noise gates, the codec's 16-bit conversion) are not ported. In the polyphonic program each voice has its own generator where a `Part` with several voices would share one.

A block-based firmware can be written as a Faust stream that gives its samples, not only its sound, if the port is checked against the firmware: the block clock and `ondemand` reproduce its control-rate work and its frozen components, the generator's jump reproduces its interleaved random draws, and the float32 constants read back from the compiled C++ reproduce what its arithmetic makes of them. The comparison found what listening would not have: a strike sample player that the port filtered where the firmware does not, 25 dB too quiet, and a reverb LFO whose frequency depended on one unit in the last place of its coefficient. And the remaining differences have causes: float rounding at the thresholds of a 16-bit memory and of a nonlinear loop, and, in double precision, the firmware's own float32 drift.

References

[1] Jean-Marie Adrien. The missing link: Modal synthesis. In Giovanni De Poli, Aldo Piccialli, and Curtis Roads, editors, *Representations of Musical Signals*, pages 269–297. MIT Press, Cambridge, MA, 1991.

[2] Perry R. Cook and Gary P. Scavone. The synthesis ToolKit (STK). In *Proceedings of the International Computer Music Conference*, Beijing, 1999.

[3] Jon Dattorro. Effect design, part 1: Reverberator and other filters. *Journal of the Audio Engineering Society*, 45(9):660–684, 1997.

[4] Georg Essl and Perry R. Cook. Banded waveguides: Towards physical modeling of bowed bar percussion instruments. In *Proceedings of the International Computer Music Conference*, Beijing, 1999.

[5] Emilie Gillet. Mutable instruments Elements: modal synthesis voice, firmware source code, 2014. `https://github.com/pichenettes/eurorack`, directory `elements`, MIT license; commit `f4cfa76`, `stmlib e3bd7c9`.

[6] David Goldberg. What every computer scientist should know about floating-point arithmetic. *ACM Computing Surveys*, 23(1):5–48, 1991.

[7] GRAME-CNCM. faust-rs: the Faust compiler in Rust, 2026. Branch `main-dev`, commit `172305ea`, 22 October 2026.

[8] David A. Jaffe and Julius O. Smith. Extensions of the Karplus-Strong plucked-string algorithm. *Computer Music Journal*, 7(2):56–69, 1983.

[9] Kevin Karplus and Alex Strong. Digital synthesis of plucked-string and drum timbres. *Computer Music Journal*, 7(2):43–55, 1983.

[10] Stéphane Letz. From prepared notes to a stream: Porting a physically modelled piano from C++ to Faust, checked against the C++ model itself, 2026. faust-coopiano, technical report, GRAME.

[11] Michael E. McIntyre, Robert T. Schumacher, and James Woodhouse. On the oscillations of musical instruments. *Journal of the Acoustical Society of America*, 74(5):1325–1345, 1983.

[12] Yann Orlarey, Dominique Fober, and Stéphane Letz. FAUST: an efficient functional approach to DSP programming. In *New Computational Paradigms for Computer Music*, pages 65–78. Delatour, Paris, 2009.

[13] William H. Press, Saul A. Teukolsky, William T. Vetterling, and Brian P. Flannery. *Numerical Recipes in C: The Art of Scientific Computing*. Cambridge University Press, 2nd edition, 1992.

[14] Julius O. Smith. Physical modeling using digital waveguides. *Computer Music Journal*, 16(4):74–91, 1992.

[15] Vadim Zavalishin. *The Art of VA Filter Design*. Native Instruments, 2012.